

Machine Learning

Lecture 05

Min-Kuan Chang

minkuanc@nchu.edu.tw

EE, College of EECS

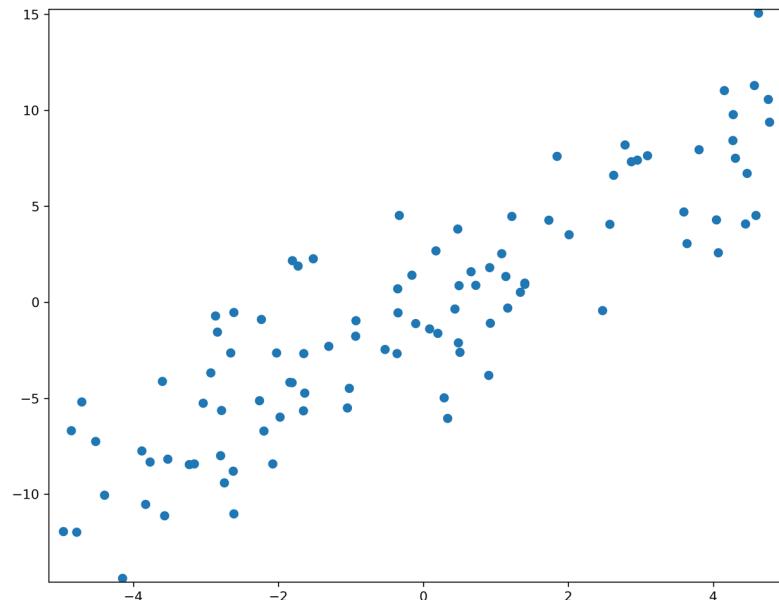
Implementation Using Scikit-Learn

Topic 01

Linear Regression

Step 01: Dataset Preparation

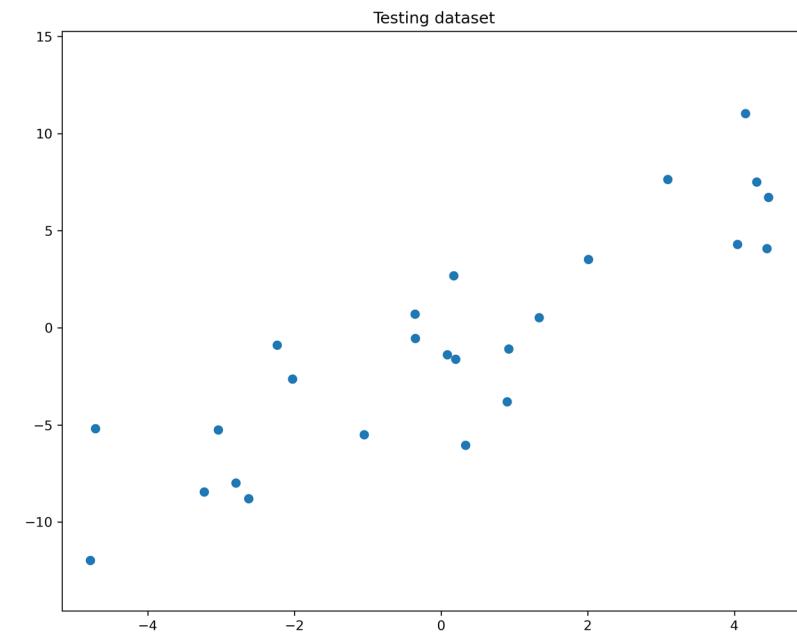
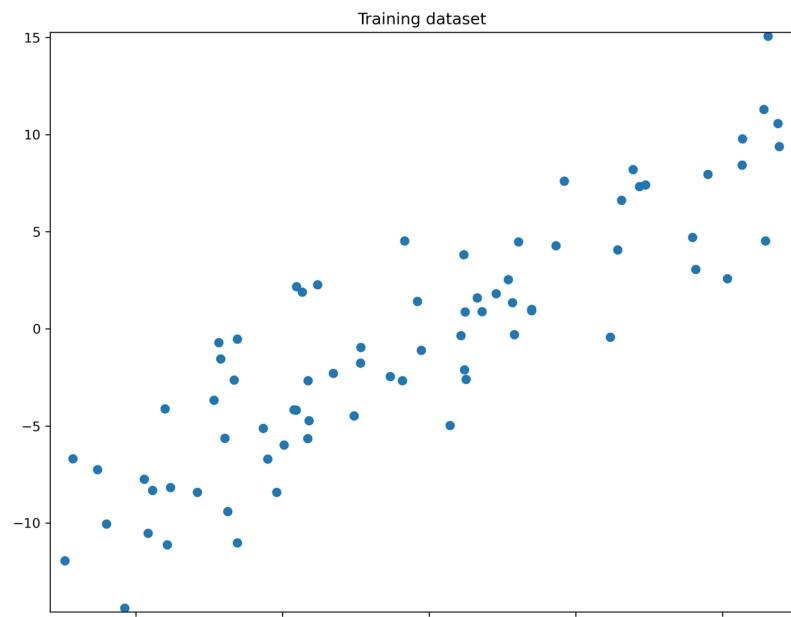
```
import numpy as np  
X = np.random.uniform(-5,5,100)  
y = 2*X + np.random.normal(0,3,100)
```



We are seeking a linear model
 $\hat{y} = w_0 + w_1 x$ to fit the data

Step 02: Training/Testing Dataset

```
from sklearn.model_selection import train_test_split  
X=X.reshape(-1,1)  
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=10)
```



Step 03: Build/Train ML Model

- ‘LinearRegression’ is responsible for building a linear regression model
- The object of class ‘LinearRegression’ is declared before training

```
from sklearn.linear_model import LinearRegression  
LinearRegression_Model = LinearRegression()
```

- no particular arguments are needed
- To train this model, we simply use the class method ‘fit’ to determine the bias and weights of the model

```
LinearRegression_Model.fit(X_train, y_train)
```

Step 04: Test the ML model

- We can use the method ‘predict’ to predict the output of a new data

```
LinearRegression_Model.predict([[3]])
```

- the output is

```
array([5.77166998])
```

- We can use the method ‘score’ to understand the performance of the model

```
LinearRegression_Model.score(X_test, y_test)
```

- the output is

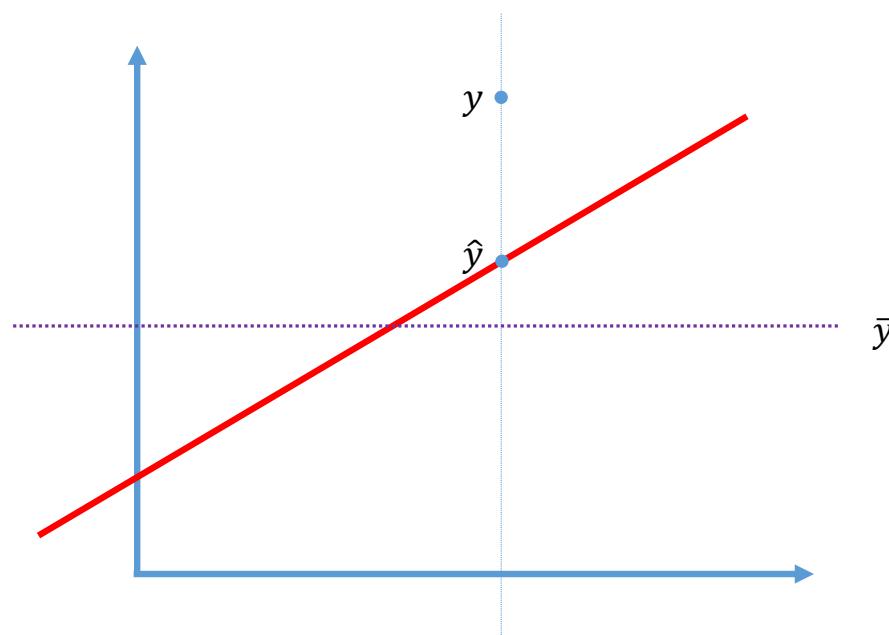
```
0.734295740699449
```

字幕預留位置

Step 04: Test the ML model

- The method ‘score’ returns the coefficient of determination or R^2 score

$$R^2 = 1 - \frac{\sum_{i=1}^N (y - \hat{y})^2}{\sum_{i=1}^N (y - \bar{y})^2}$$



Discussion

- LinearRegression_Model has many useful attributes
 - the information of the fitting function

```
LinearRegression_Model.intercept_
```

- the bias term, w_0 , of the linear model
- the output is

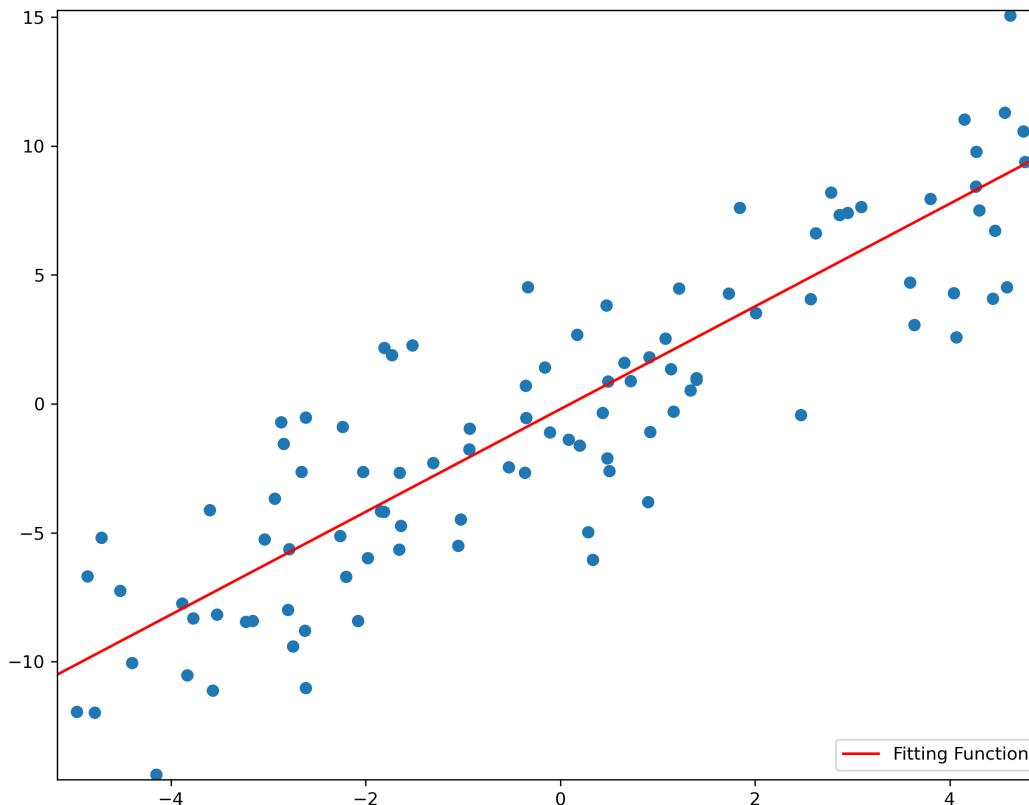
```
-0.2047650511796517
```

```
LinearRegression_Model.coef_
```

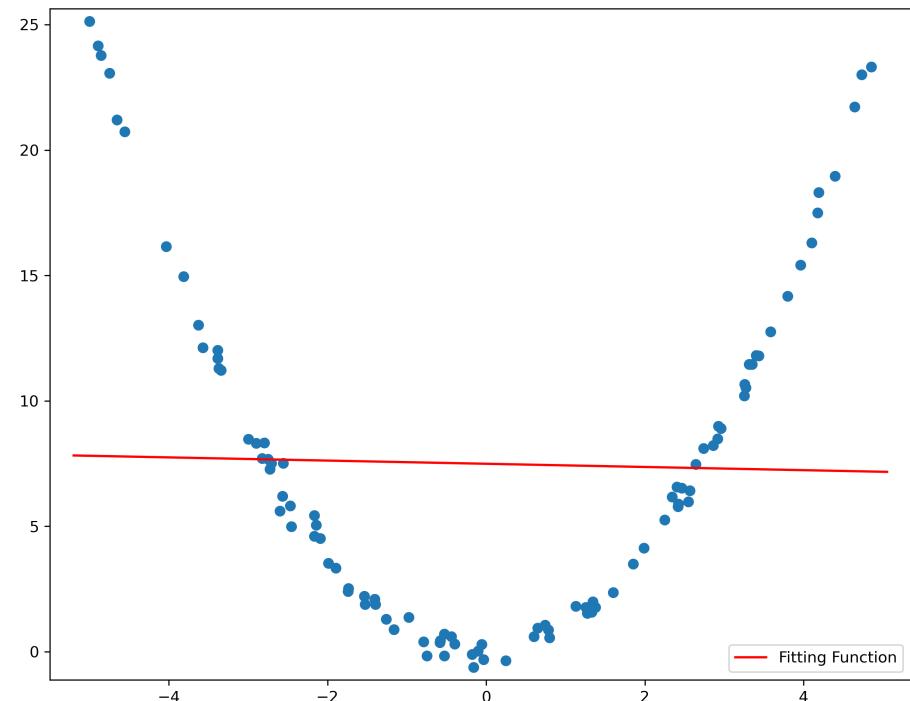
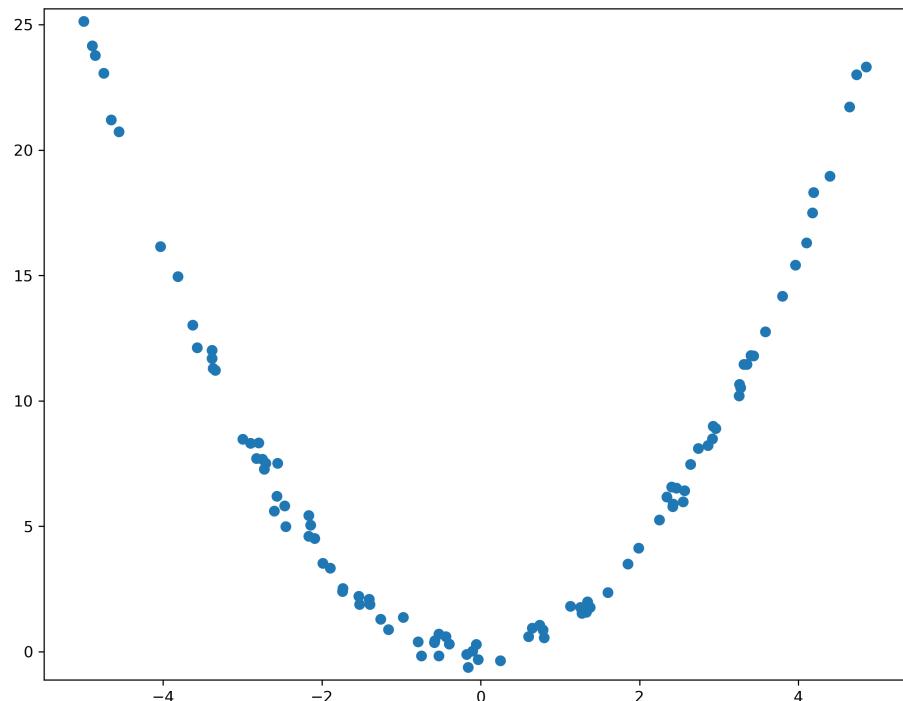
- record the weights of the linear model
- the output is

```
array([1.99214501])
```

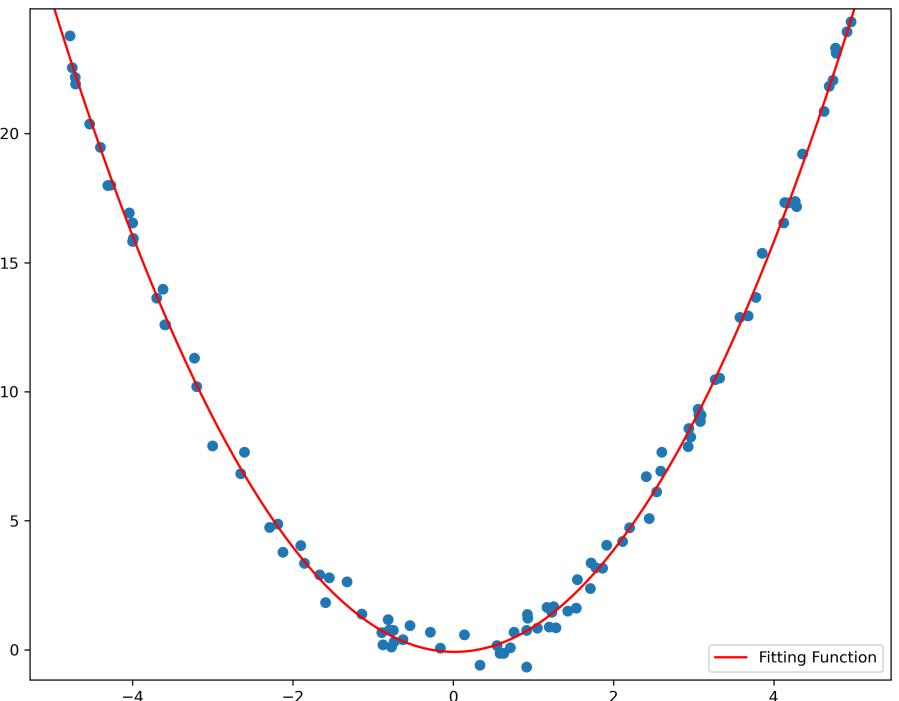
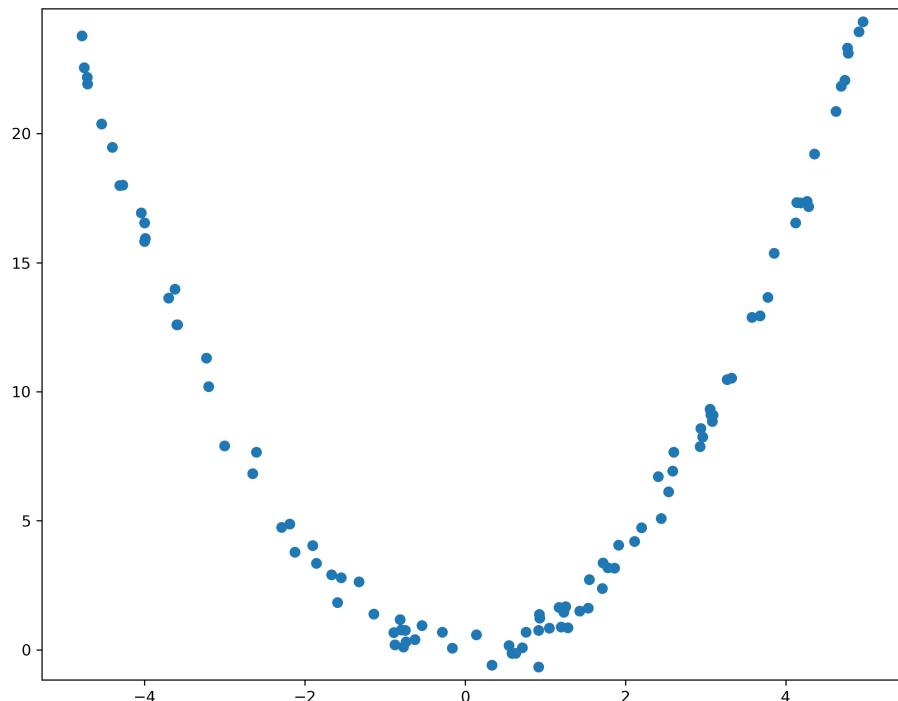
Discussion



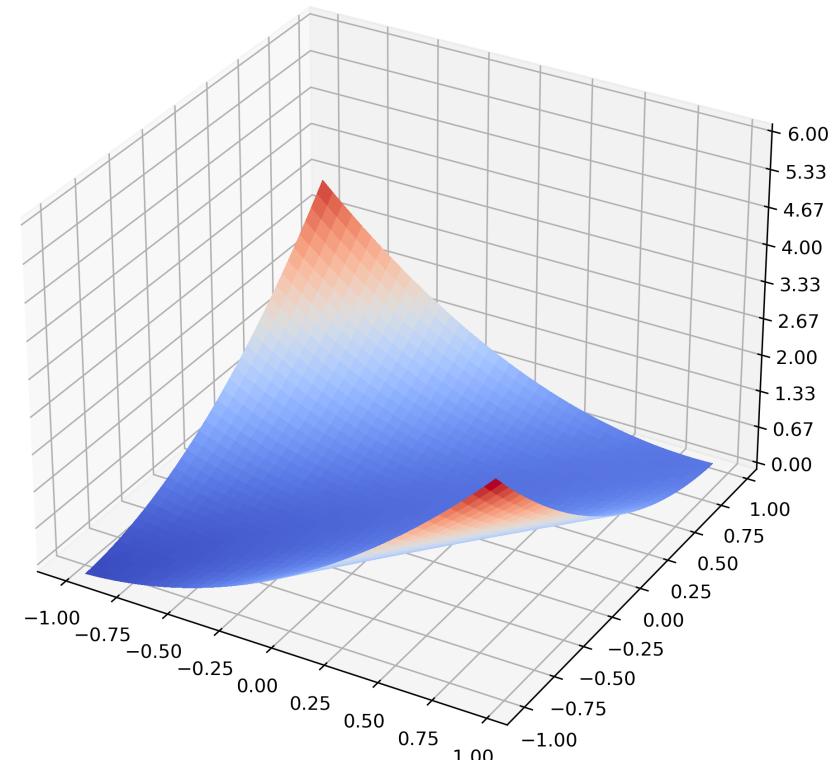
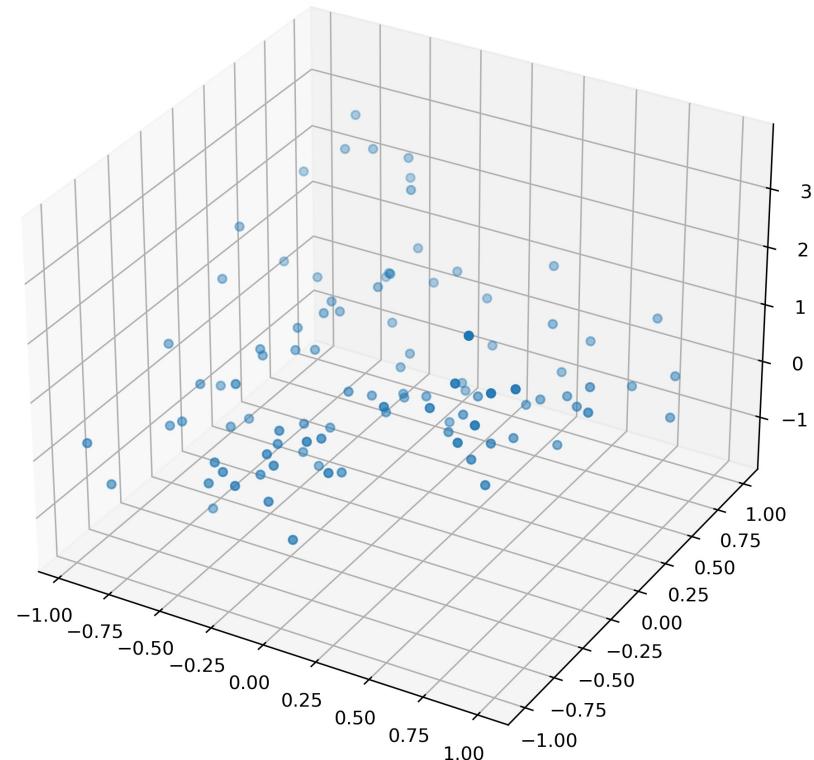
Discussion



Discussion



Discussion



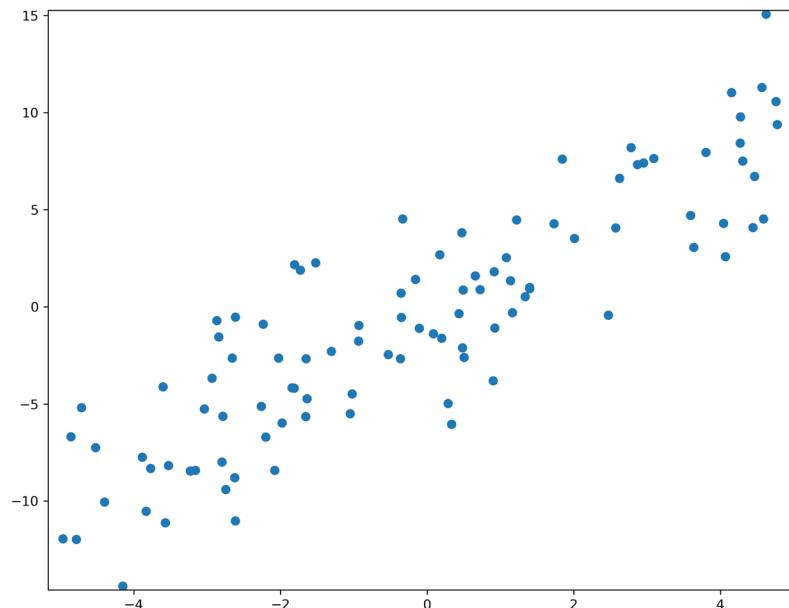
Implementation Using Scikit-Learn

Topic 02

Regularized Linear Regression - Ridge

Step 01: Dataset Preparation

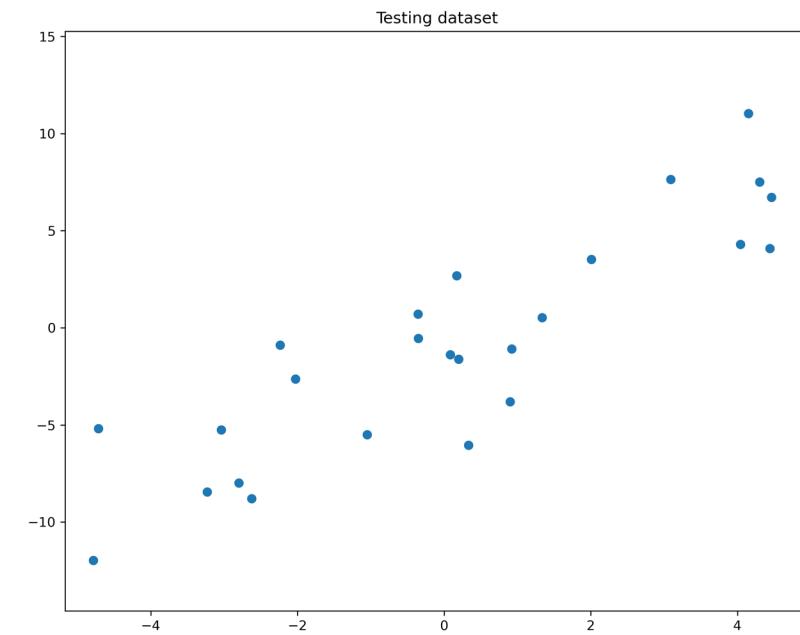
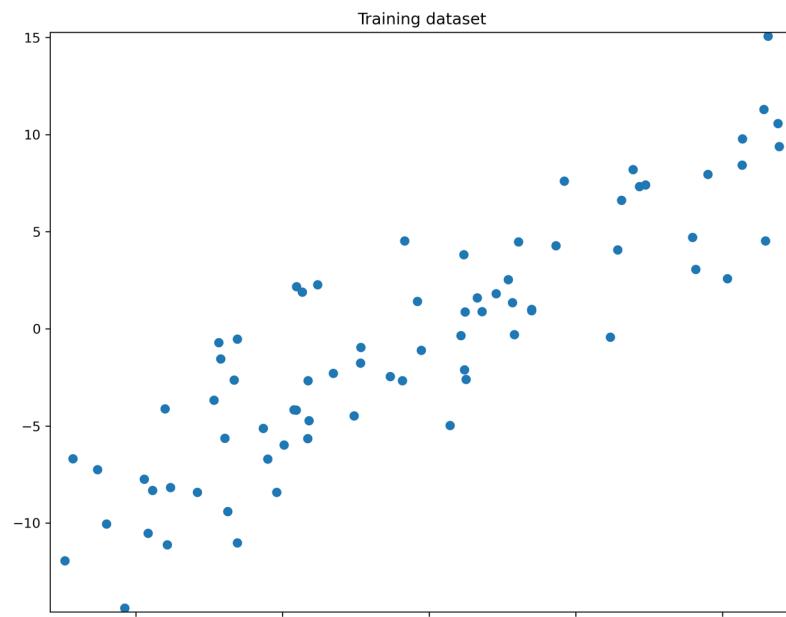
```
import numpy as np  
X = np.random.uniform(-5,5,100)  
y = 2*X + np.random.normal(0,3,100)
```



We are seeking for a Ridge model
 $\hat{y} = w_0 + w_1x$ to fit the data

Step 02: Training/Testing Dataset

```
from sklearn.model_selection import train_test_split  
X=X.reshape(-1,1)  
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=10)
```



Step 03: Build/Train ML Model

- The class ‘Ridge’ is utilized to build a Ridge model
- We simply declare an object of ‘Ridge’ and initialize this model

```
from sklearn.linear_model import Ridge  
Ridge_Model = Ridge(alpha = 1, max_iter = 1000, tol = 0.001)
```

- ‘alpha’ controls the strength of the regularization
 - ‘max_iter’ determines the maximum number of iteration to find the weights of the models
 - ‘tol’ is the precision of the solution
- The training of the model is to call the ‘fit’ method

```
Ridge_Model.fit(X_train, y_train)
```

Step 04: Test the ML model

- When a new input comes, we can use the ‘predict’ method to predict its possible output

```
Ridge_Model.predict([[5]])
```

- its output is

```
array([8.91188908])
```

- The ‘score’ method is to evaluate how fit the model is

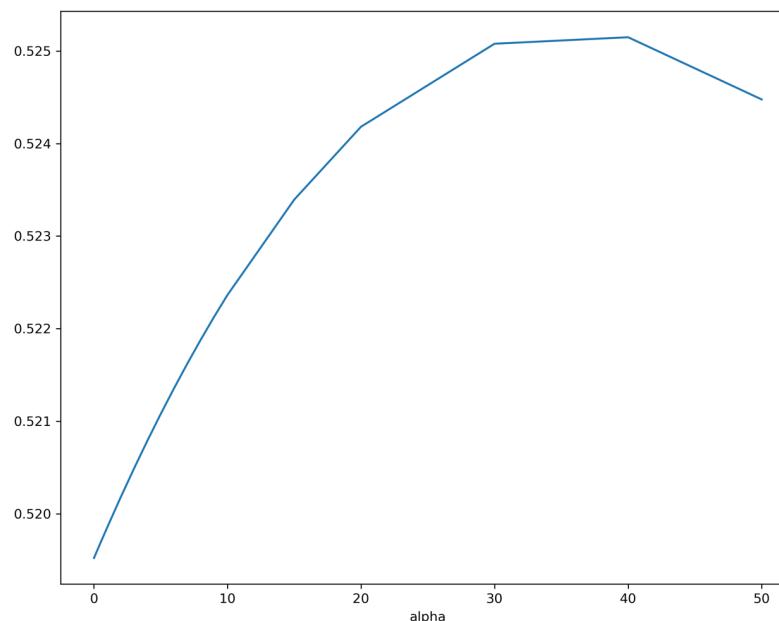
```
Ridge_Model.score(X_test, y_test)
```

- its output is

```
0.5198518682703732
```

Discussion

- Since the regularization plays a part in Ridge, we need to access how the strength of regularization affects the R^2 score



Discussion

- The bias and the weight of the fitting function
 - bias: w_0

```
Ridge_Model.intercept_
```

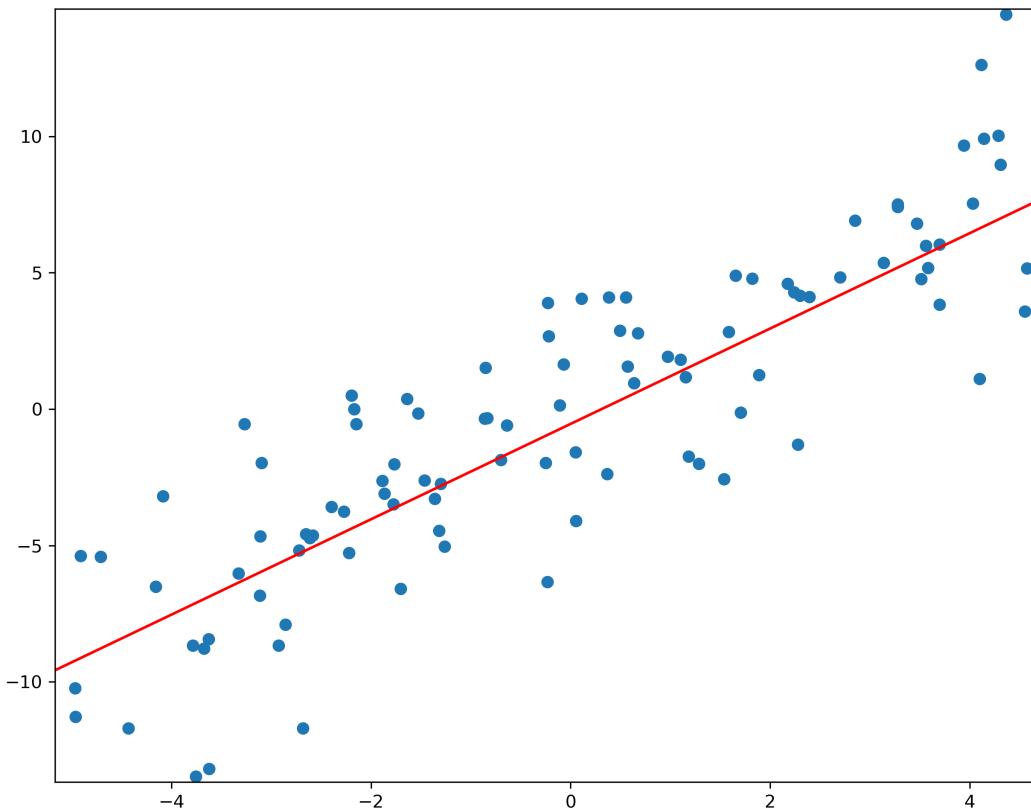
```
-0.5441973828355591
```

- the weight: w_1

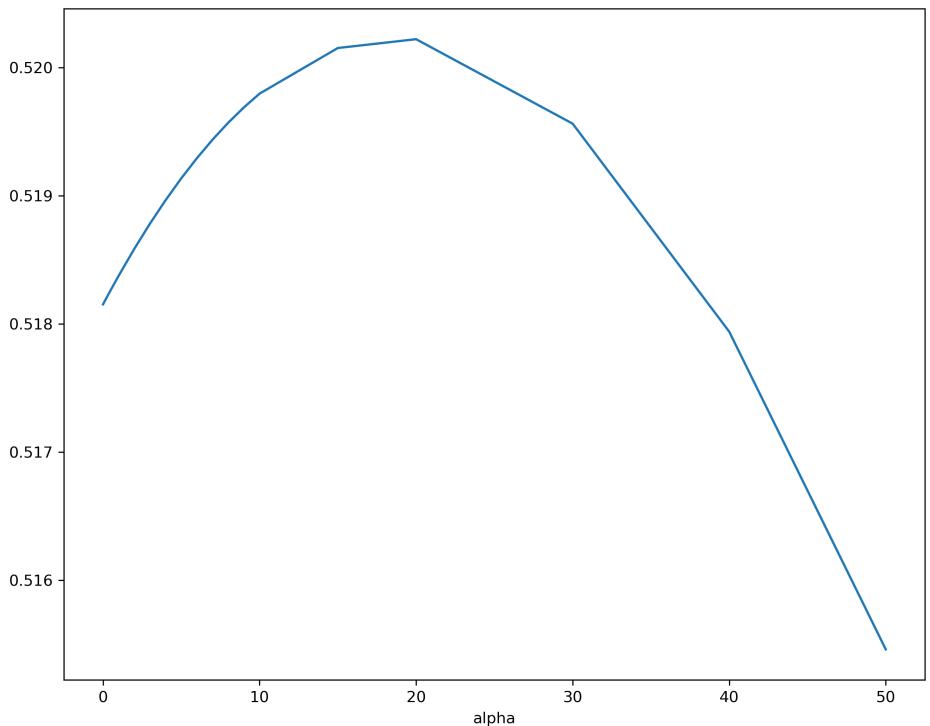
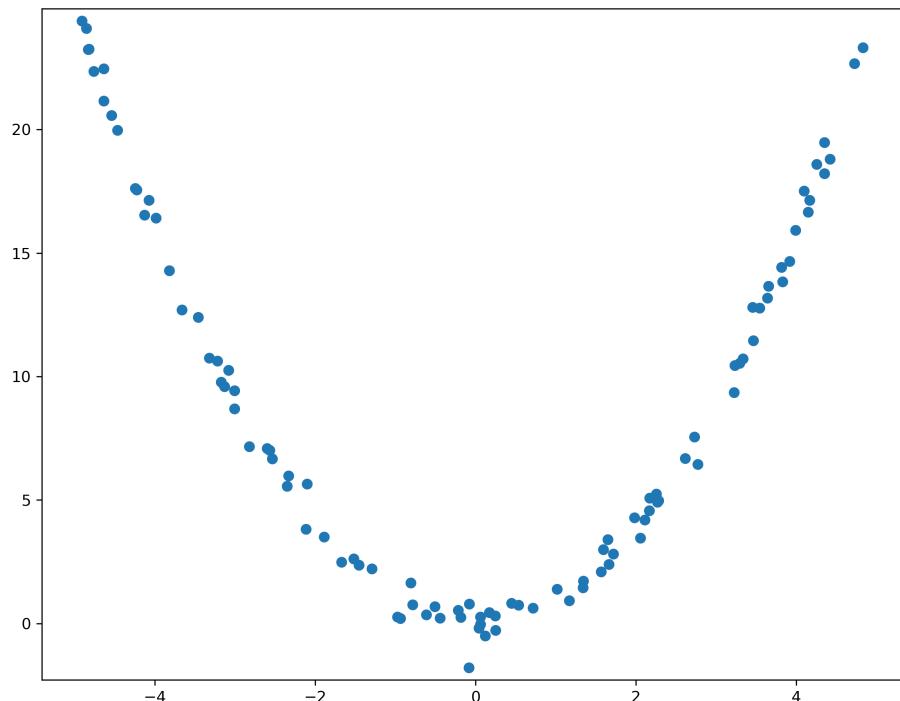
```
Ridge_Model.coef_
```

```
array([1.74817381])
```

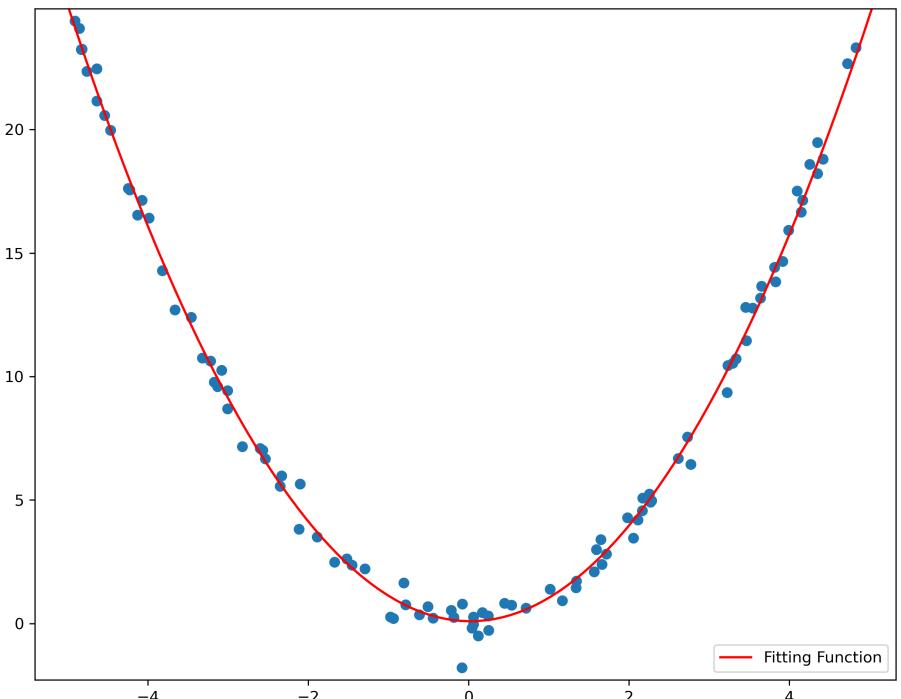
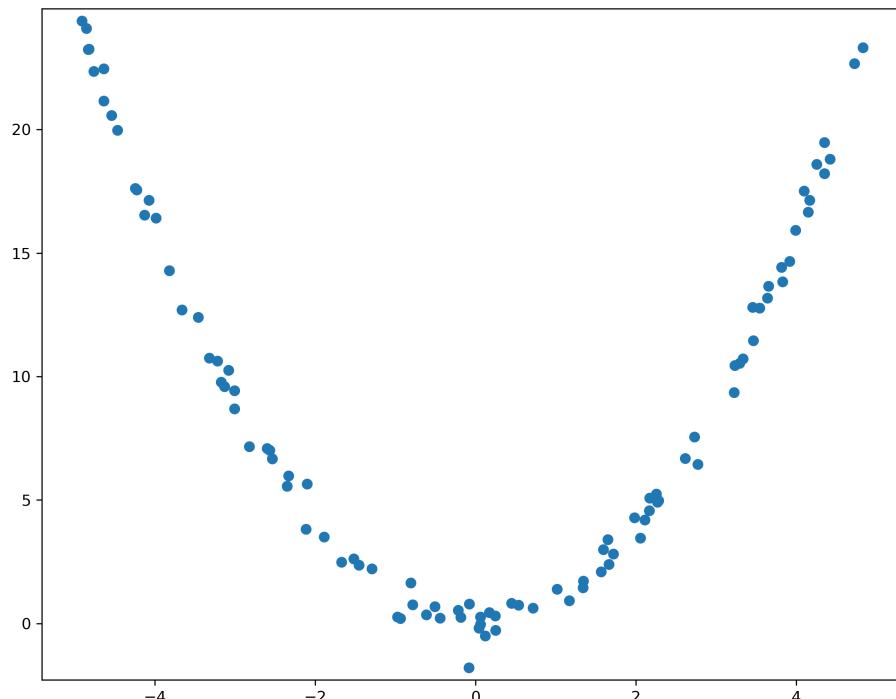
Discussion



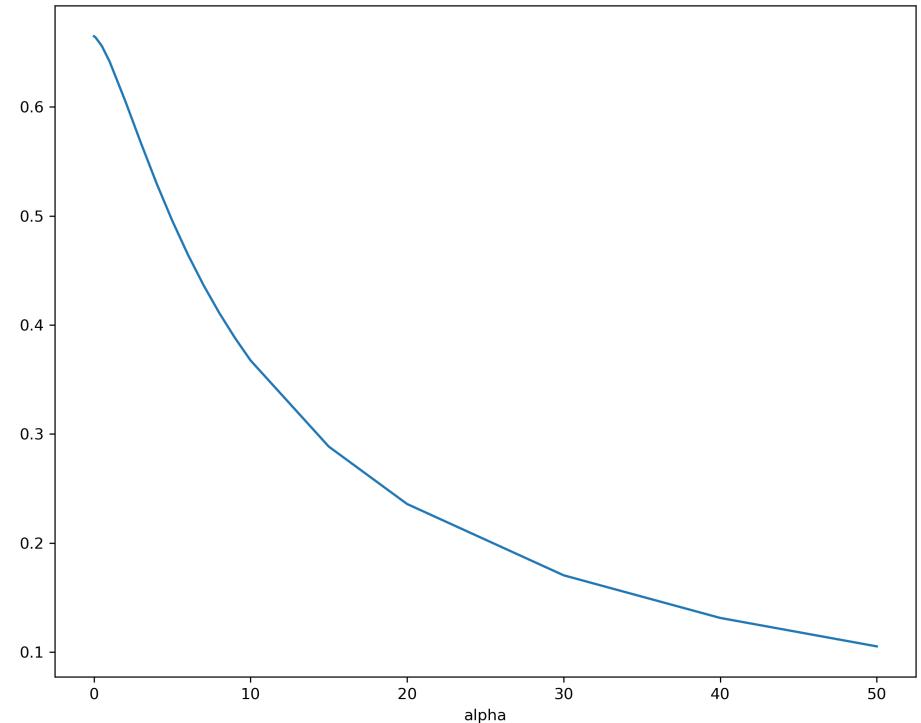
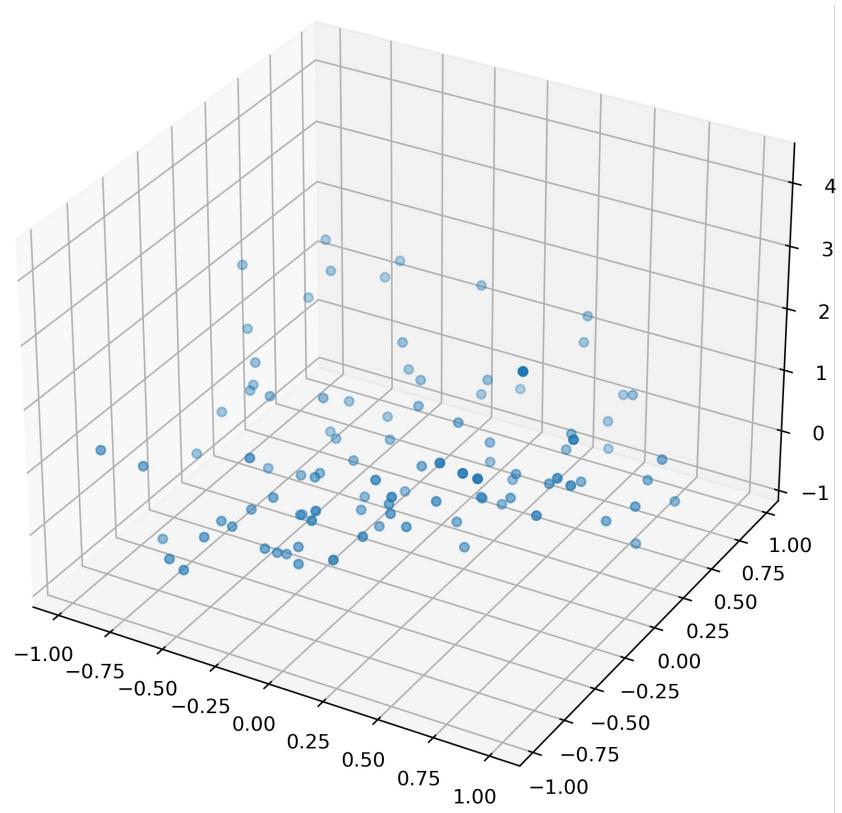
Discussion



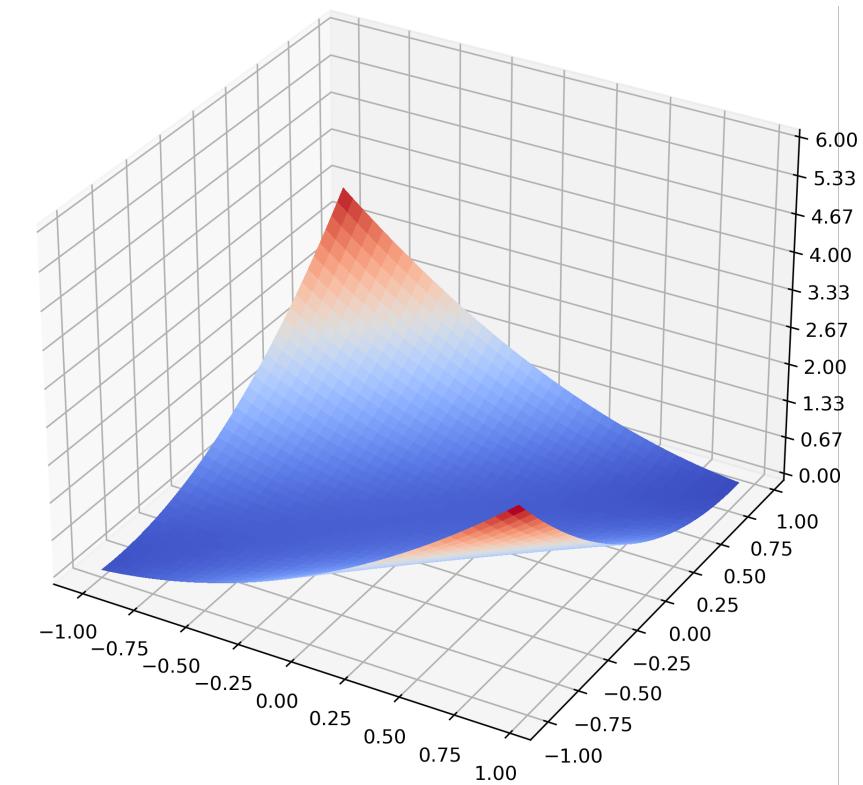
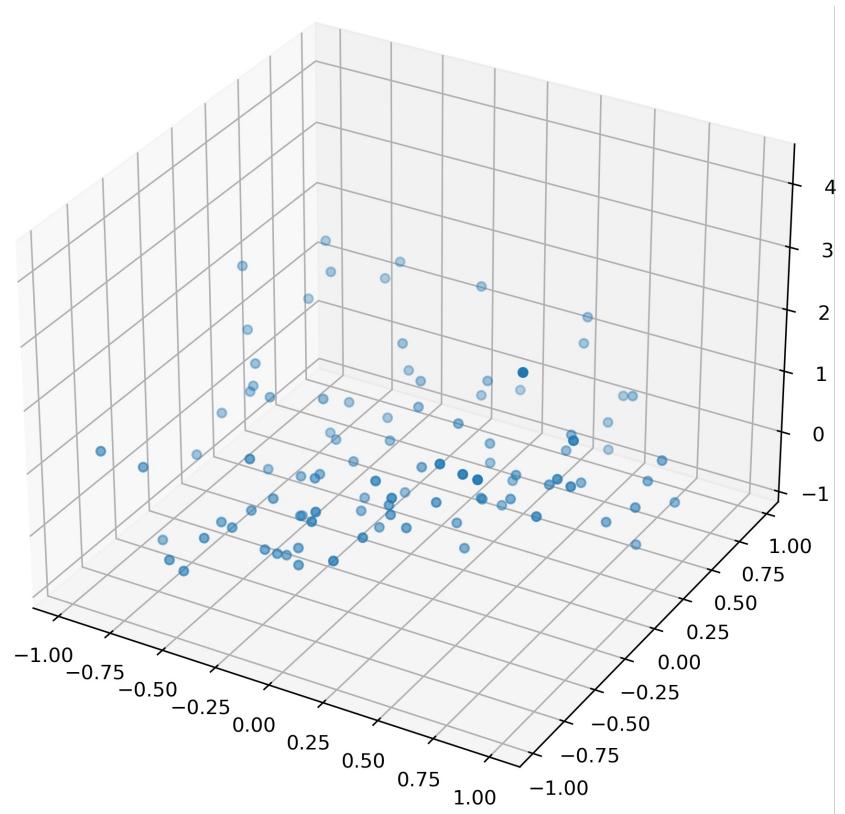
Discussion



Discussion



Discussion



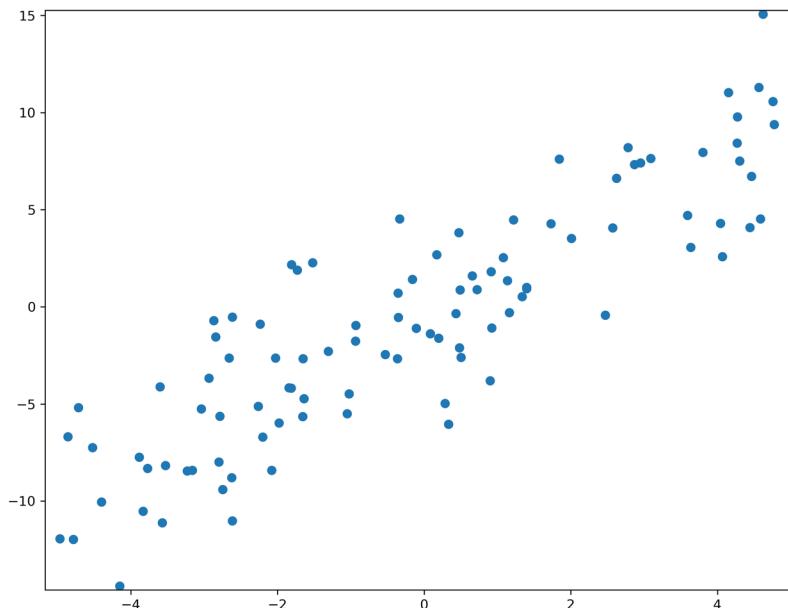
Implementation Using Scikit-Learn

Topic 03

Regularized Linear Regression - Lasso

Step 01: Dataset Preparation

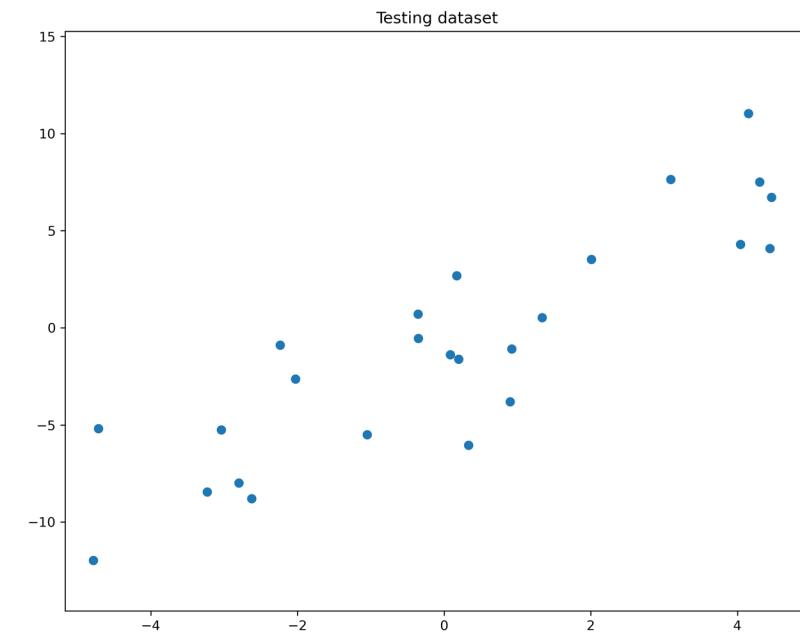
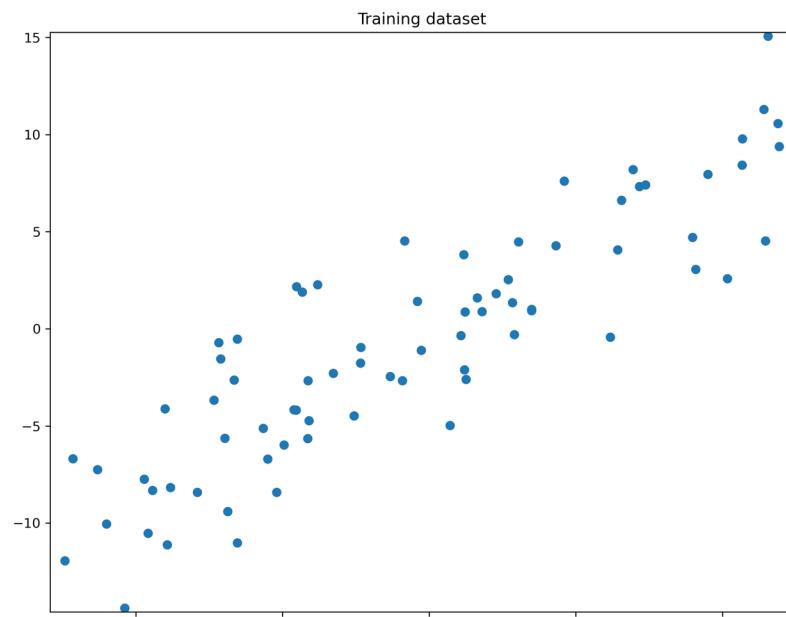
```
import numpy as np  
X = np.random.uniform(-5,5,100)  
y = 2*X + np.random.normal(0,3,100)
```



We are seeking for a Lasso model
 $\hat{y} = w_0 + w_1 x$ to fit the data

Step 02: Training/Testing Dataset

```
from sklearn.model_selection import train_test_split  
X=X.reshape(-1,1)  
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=10)
```



Step 03: Build/Train ML Model

- The class ‘Lasso’ is utilized to build a Lasso model
- An object of ‘Lasso’ is declared to initialize this model

```
from sklearn.linear_model import Lasso  
Lasso_Model = Lasso(alpha = 1, max_iter = 1000, tol = 0.001)
```

- ‘alpha’ controls the strength of the regularization
 - ‘max_iter’ determines the maximum number of iteration to find the weights of the models
 - ‘tol’ is the precision of the solution
- The training of the model is to call the ‘fit’ method

```
Lasso_Model.fit(X_train, y_train)
```

Step 04: Test the ML model

- When a new input comes, we can use the ‘predict’ method to predict its possible output

```
Lasso_Model.predict([[5]])
```

- its output is

```
array([12.54081639])
```

- The ‘score’ method is to evaluate how fit the model is

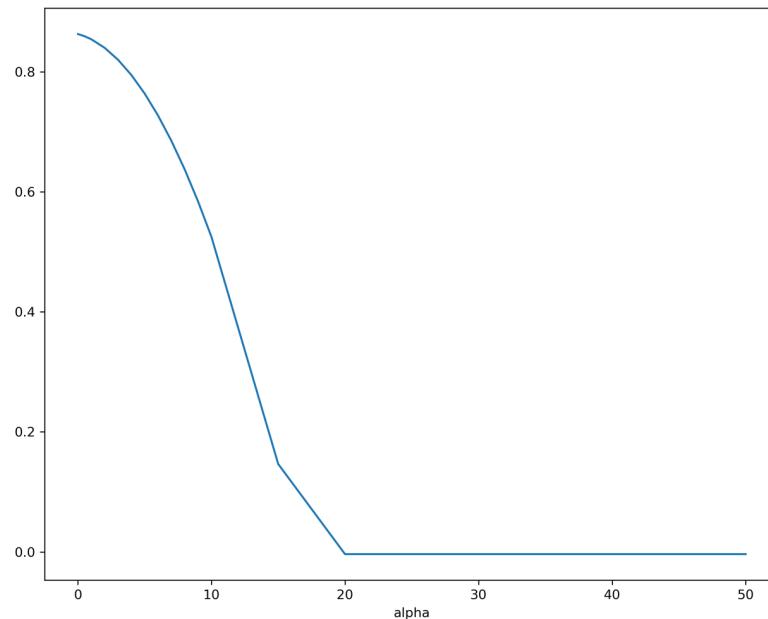
```
Lasso_Model.score(X_test, y_test)
```

- its output is

```
0.8542612376965604
```

Discussion

- Since the regularization plays a part in Lasso, we need to access how the strength of regularization affects the R^2 score



Discussion

- The bias and the weight of the fitting function
 - bias: w_0

```
Lasso_Model.intercept_
```

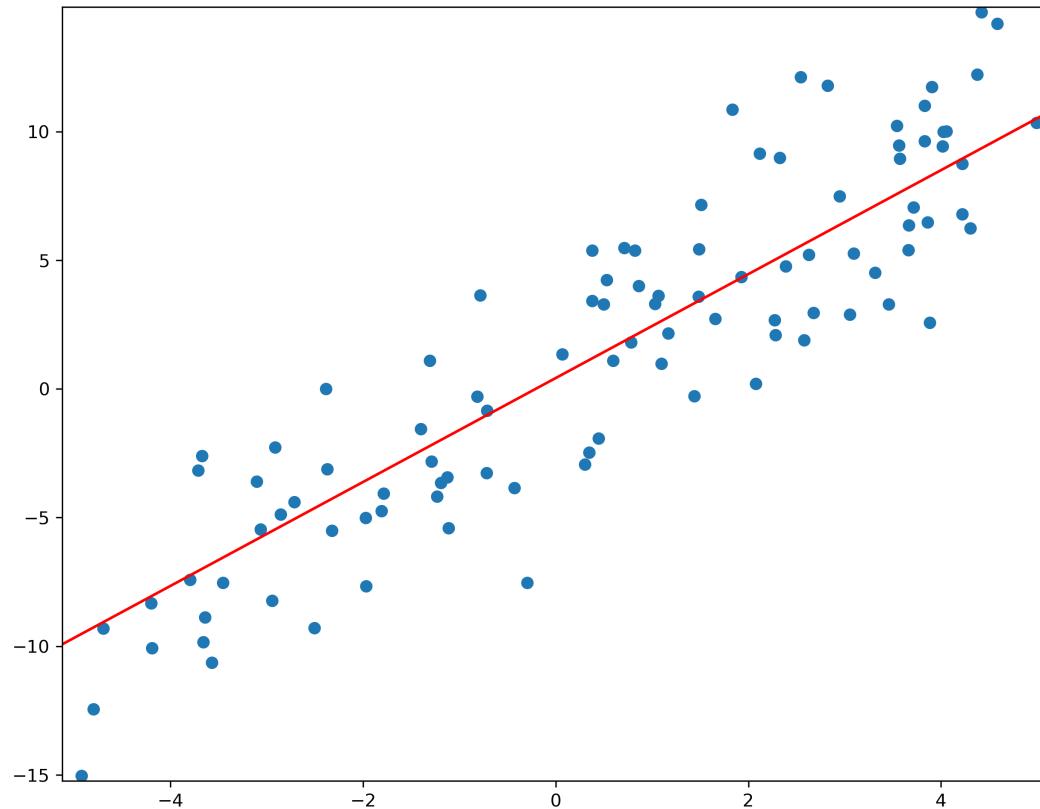
```
0.4235520049058268
```

- the weight: w_1

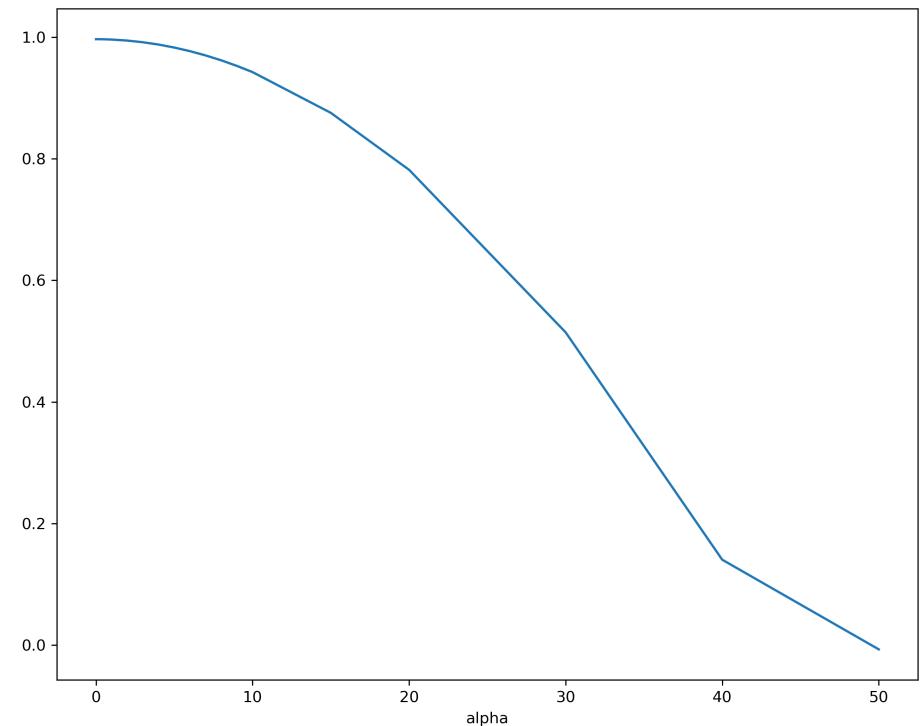
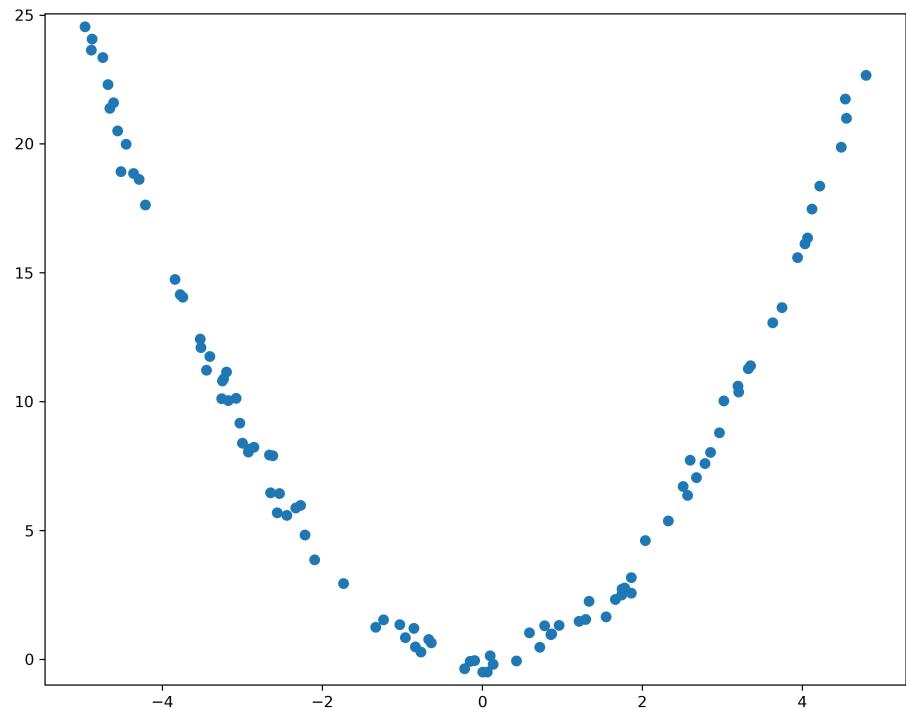
```
Lasso_Model.coef_
```

```
array([2.01954406])
```

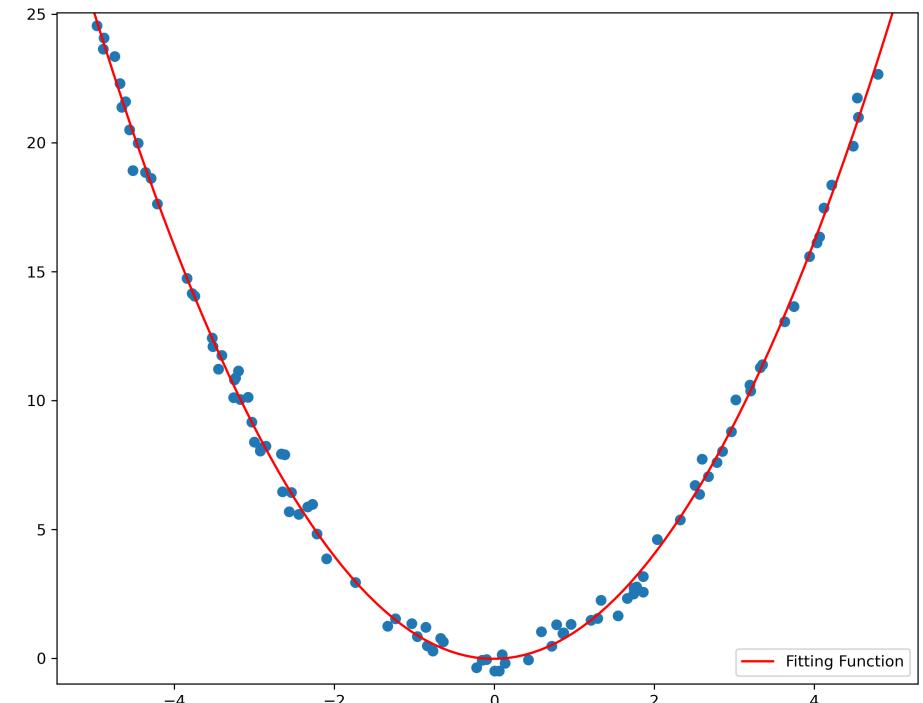
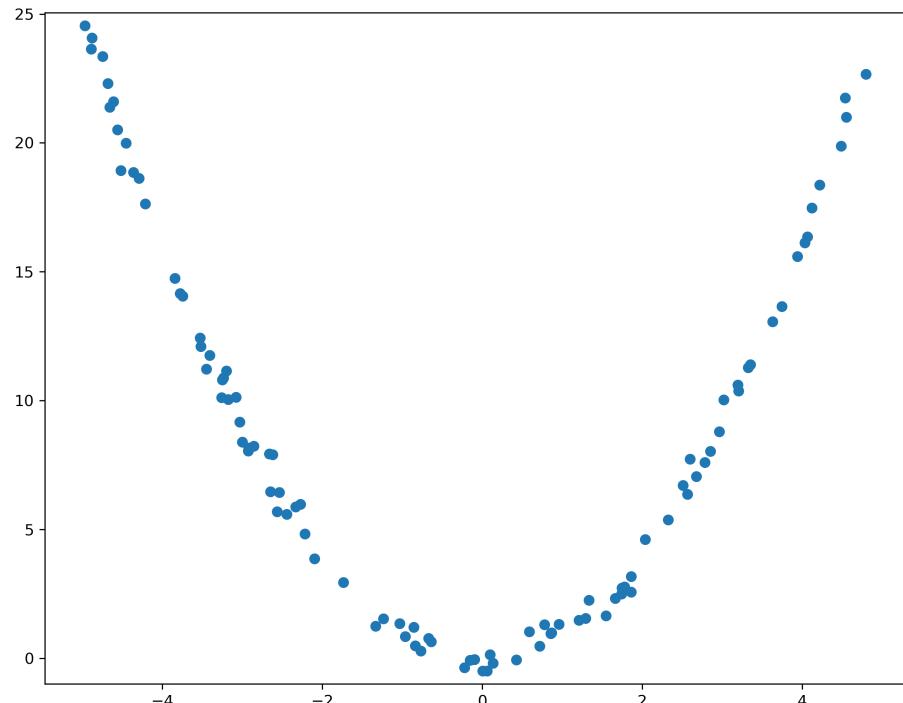
Discussion



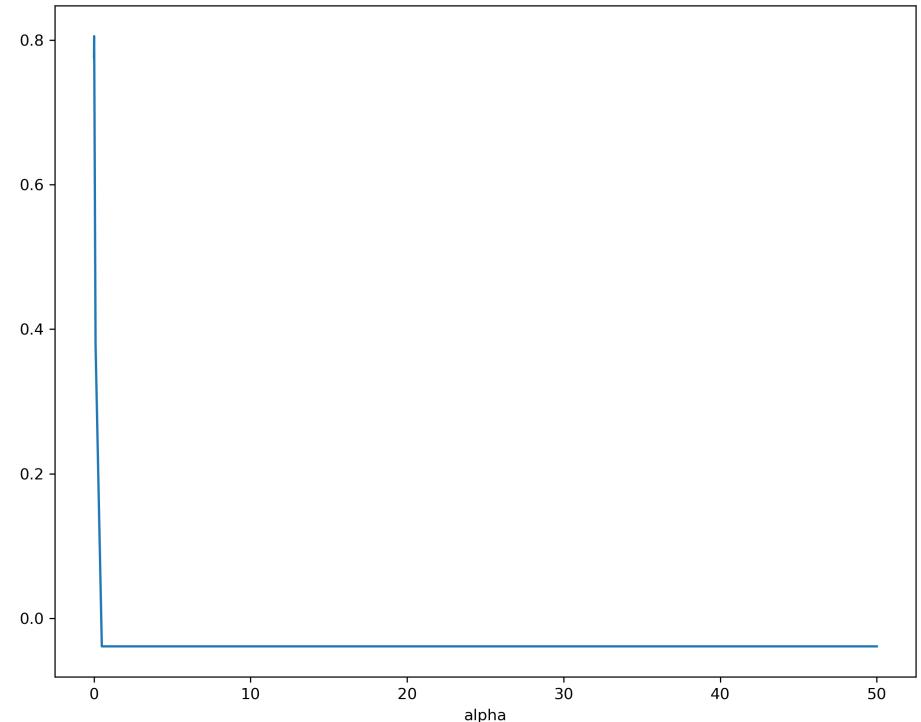
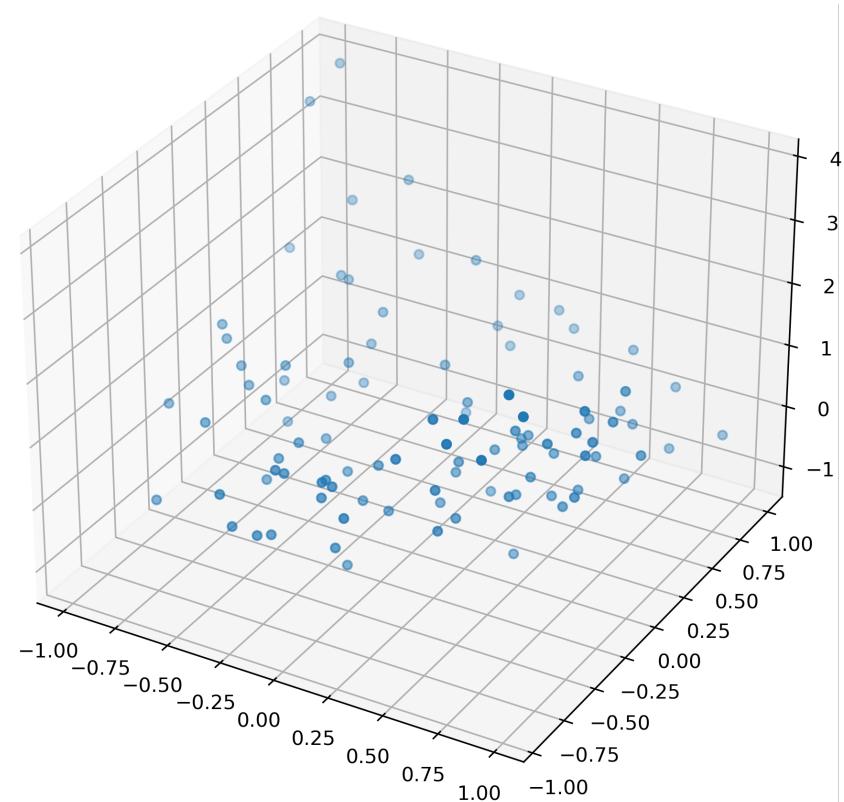
Discussion



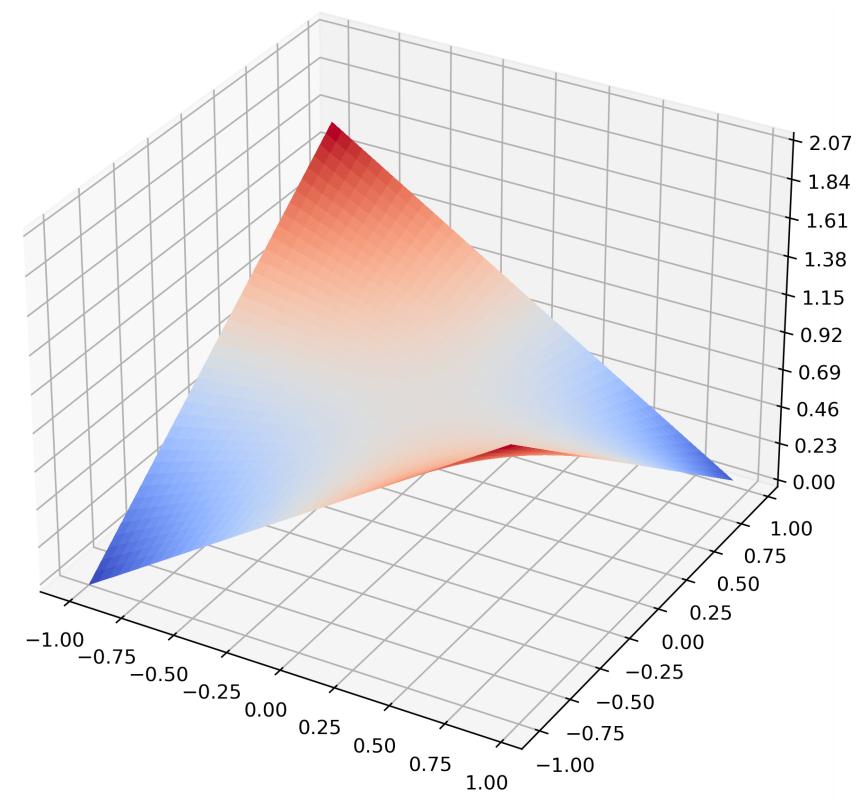
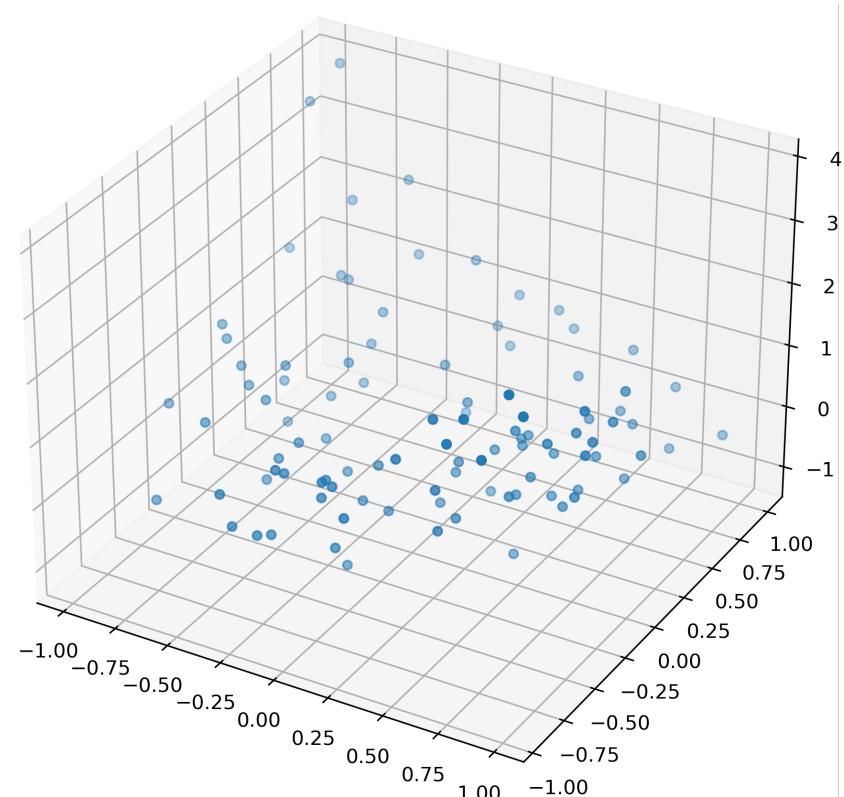
Discussion



Discussion



Discussion



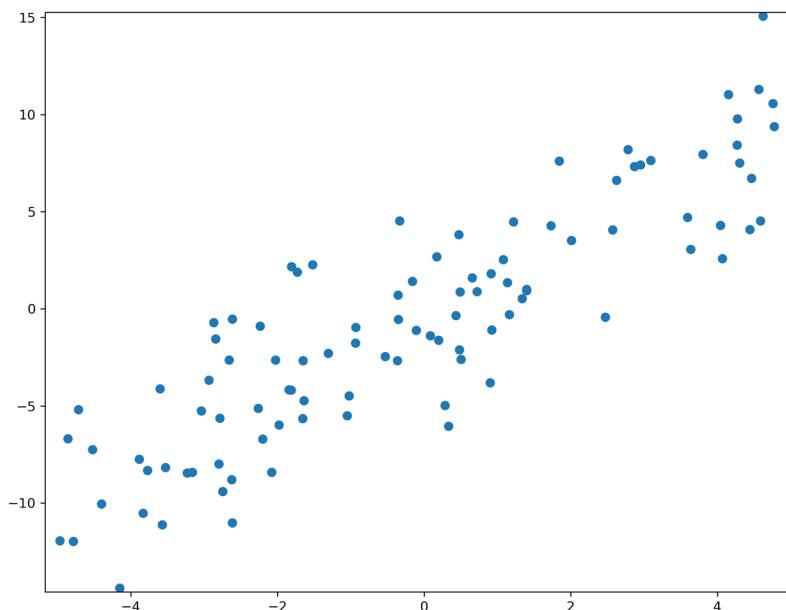
Implementation Using Scikit-Learn

Topic 04

Bayesian Regression

Step 01: Dataset Preparation

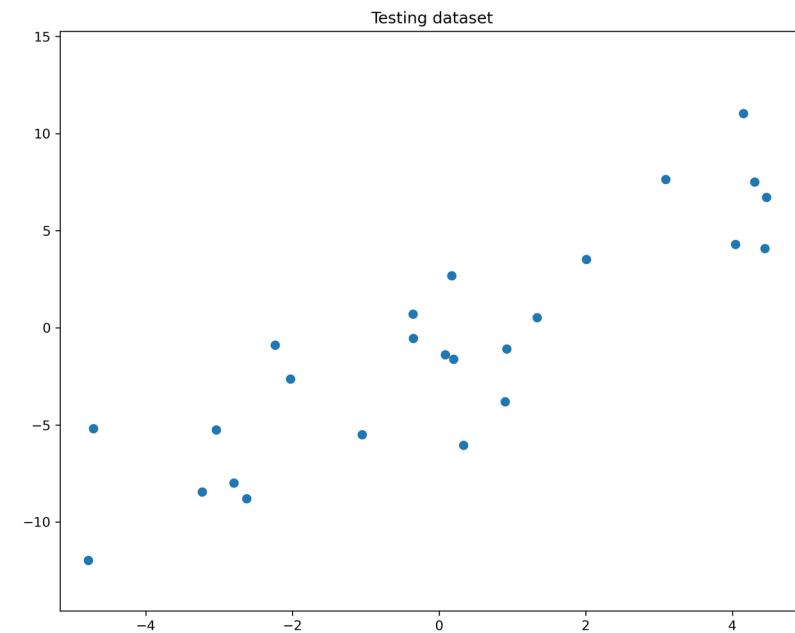
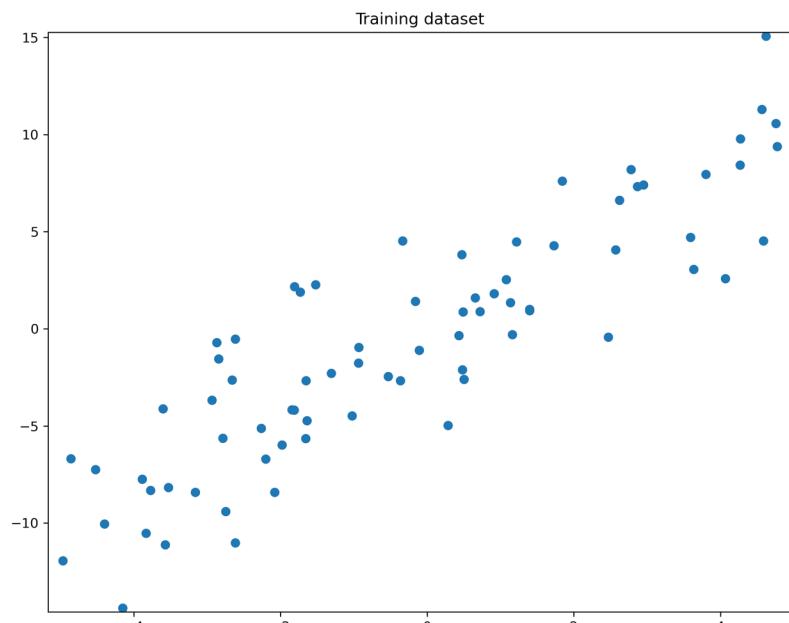
```
import numpy as np  
X = np.random.uniform(-5,5,100)  
y = 2*X + np.random.normal(0,3,100)
```



We are seeking for a Bayesian regression model
 $\hat{y} = w_0 + w_1x$ to fit the data

Step 02: Training/Testing Dataset

```
from sklearn.model_selection import train_test_split  
X=X.reshape(-1,1)  
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=10)
```



Step 03: Build/Train ML Model

- ‘BayesianRidge’ helps build the Bayesian regression in Scikit-learn
- Two prior distributions are assumed in ‘BayesianRidge’
 - the prior distribution of the variance of observation follows the gamma distribution with parameters ‘alpha_1’ and ‘alpha_2’
 - the prior distribution of the variance of weights follows the gamma distribution with parameters ‘lambda_1’ and ‘lambda_2’

Step 03: Build/Train ML Model

```
from sklearn.linear_model import BayesianRidge  
BayesianRidge_Model = BayesianRidge(  
    n_iter=300, tol=0.001,  
    alpha_1=1e-06, alpha_2=1e-06,  
    lambda_1=1e-06, lambda_2=1e-06  
)
```

- ‘n_iter’ determines the maximum number of iteration to find the weights of the models
- ‘tol’ is the precision of the solution
- The training of the model is to call the ‘fit’ method

```
BayesianRidge_Model.fit(X_train, y_train)
```

Step 04: Test the ML model

- When a new input comes, we can use the ‘predict’ method to predict its possible output

```
BayesianRidge_Model.predict([[1.6]])
```

- its output is

```
array([3.27289468])
```

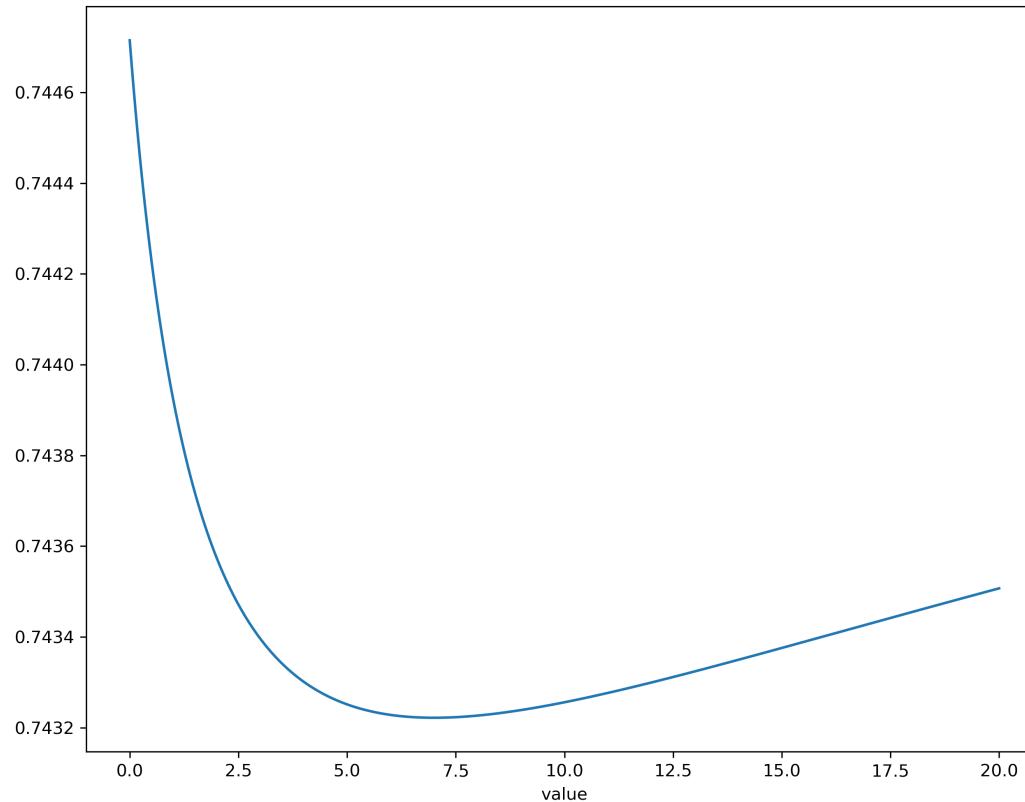
- The ‘score’ method is to evaluate how fit the model is

```
BayesianRidge_Model.score(X_test, y_test)
```

- its output is

```
0.7440097954763434
```

Discussion



Discussion

- The bias and the weight of the fitting function
 - bias: w_0

```
BayesianRidge_Model.intercept_
```

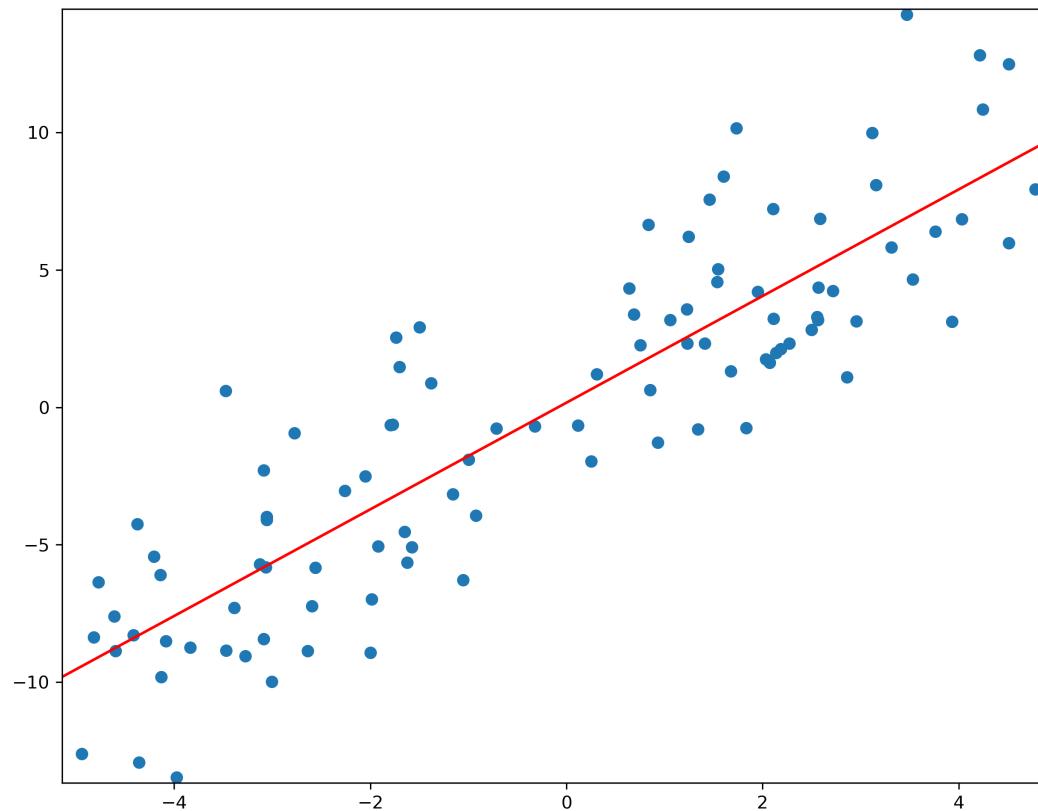
```
0.1670676923915525
```

- the weight: w_1

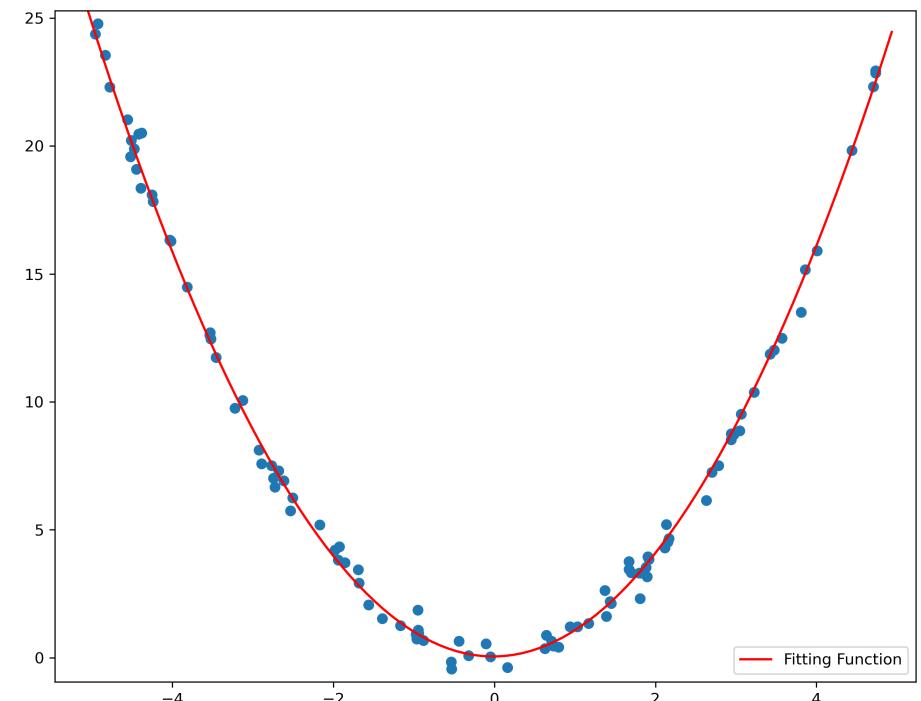
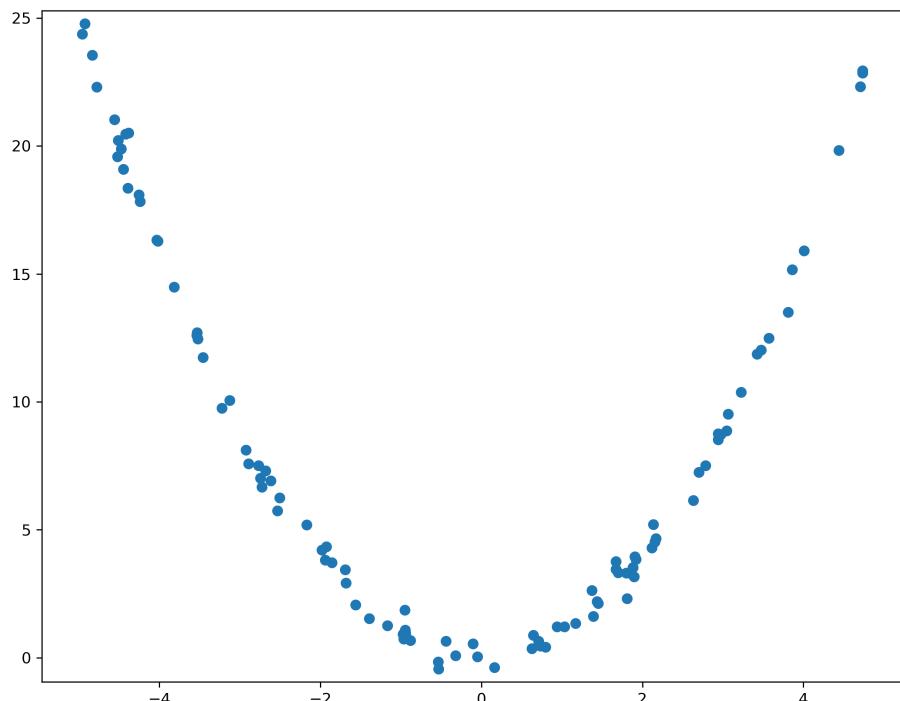
```
BayesianRidge_Model.coef_
```

```
array([1.94114187])
```

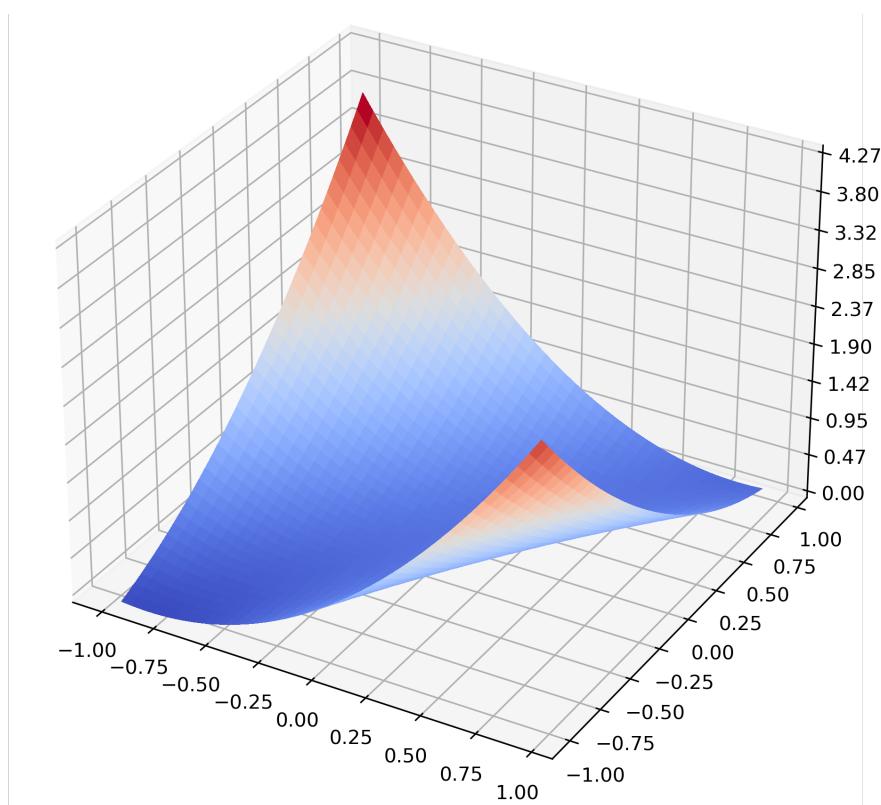
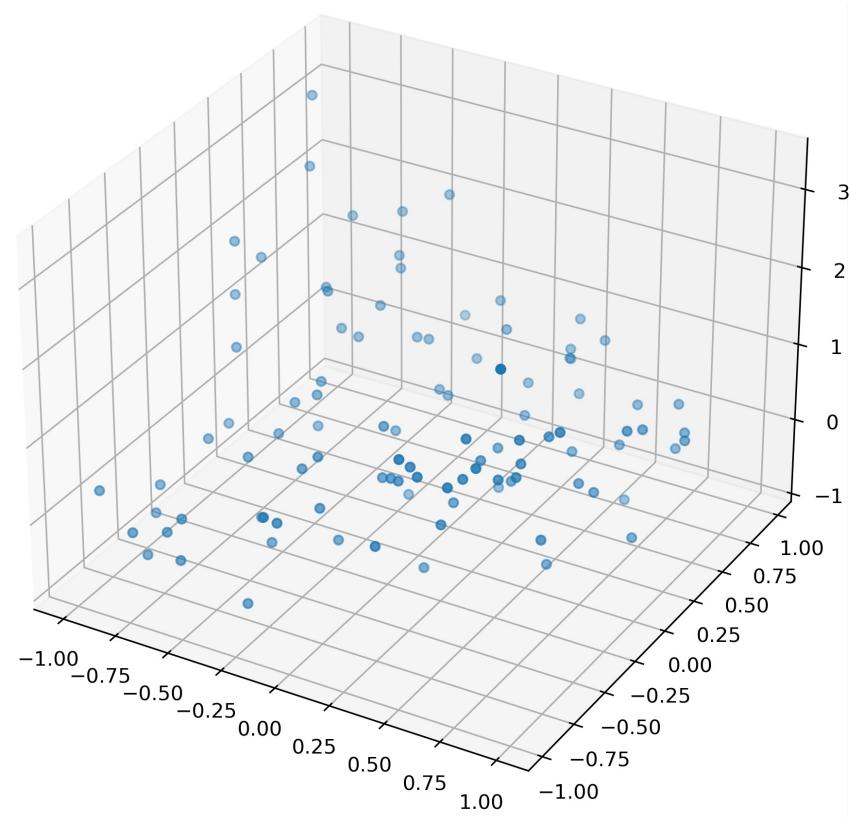
Discussion



Discussion



Discussion



Q&A

National Chung Hsing University
Wireless Multimedia and Communication Lab.